

Module 12: Lab

Elita Danilyuk

1. Rigorous 2PL with timestamps used for deadlock prevention (using Wait-Die policy).

- S1:

T1 acquires shared lock on X

for an exclusive lock on X, T2 will be aborted when asked because it has a lower priority

T3 gets an exclusive lock on Y

when T1 asks for an exclusive lock on Y (still held by T3) since T1 has a higher priority, so T1 will be blocked waiting

T3 finishes writer, commits, then releases all locks

T1 wakes up, gets the exclusive lock on Y and finishes
T2 can now be restarted successfully

- S2: similar to the sequence of S1 above, except T2 would be able to advance a bit further before it gets aborted

2. Rigorous 2PL with deadlock detection.

(Show the waits-for graph in case of deadlock.)

- Deadlock detection allows transactions to wait and they're not aborted until a deadlock is detected.

- S1:

T1 gets shared lock on X

T2 blocks waiting for exclusive lock on X

T3 gets exclusive lock on Y

T1 blocks waiting for exclusive lock on Y

T3 finishes, commits, releases locks

T1 wakes up, gets exclusive lock on Y, finishes, releases locks on X + Y

T2 gets exclusive locks on X + Y then finishes
(No deadlock)

- S2: a deadlock occurs, T1 waits for T2, while T2 waits for T1

3. Timestamp concurrency control with buffering of reads and writes (to ensure recoverability) and the Thomas Write Rule.

- initiate timestamp of X and Y to 0

- S1:

T1: R(X) allowed $TS(T1) + WTS(X) = 0 \rightarrow TS(T1) \geq WTS(X)$, RTS(X) $\rightarrow 0$

T2: W(X) allowed $TS(T2) = 2$, RTS(X) = 1, WTS(X) = 0

$\rightarrow TS(T2) \geq RTS(X) \wedge TS(T2) \geq WTS(X)$. WTS(X) $\rightarrow 2$

T3: W(Y) allowed $TS(T3) = 3$, RTS(Y) = 0, WTS(Y) = 2

$\rightarrow TS(T2) \geq RTS(X) \wedge TS(T2) \geq WTS(X)$. WTS(X) $\rightarrow 3$

T1: W(Y) ignored $TS(T1) = 1$, RTS(Y) = 0, WTS(Y) = 3

$\rightarrow TS(T2) \geq RTS(X) \wedge TS(T1) < WTS(Y)$

- S2: same as sequence S1 above

4. Validation (Optimistic) version control

- Both transactions will execute, read from the db and write to a private workspace. Then they get a timestamp to enter validation. The timestamp of transaction T_i is i .

- S1: T_1 gets earliest timestamp so it commits without any issue, but when T_1 validates T_2 against T_1 , one of the following must hold true:

• T_1 completes before T_2 begins

• T_1 completes before T_2 starts a write & T_1 doesn't write any db object read by T_2

• T_1 completes its read before T_2 completes its write & T_1 doesn't write any db object that's read or written by T_2

none of the hold true, so T_2 will be aborted and restarted later, so T_3 is the same as T_2

- S2: same as sequence S1 above

5. Multiversion timestamp concurrency control

- S1: T1 reads X therefore $RTS(X) = 1$

T2 can write X; $TS(T2) \geq RTS(X)$, $RTS(X) + WTS(X) \rightarrow 2$

T2 writes Y; $RTS(Y) + WTS(Y) \rightarrow 2$

T3 can write to Y; $TS(T3) < RTS(Y)$

→ T1 would need to be aborted & restarted later
with a larger timestamp

- S2: similar to the sequence of S1 above